

Problems Caused by Threads

1. Create thread-safe libraries
2. Create cancel-safe libraries
3. Adapt semantics of UNIX's single-threaded model
4. Synchronize access to shared variables

1. Thread-safe Libraries

For library to be “thread safe” means: any number of threads may be executing functions from this library simultaneously

The danger: shared data

Library functions must be **reentrant**:

- Function does not return pointer to static
- Function does not write to shared `errno`
- Function does not use globals OR function gets lock before accessing globals

Thread-safe Libraries, II

During last two decades, C library implementations have been re-written to be thread safe:

- Function does not return pointer to static; e.g., `strerror_r` replaces `strerror`
- Each thread has its own `errno` located in per-thread storage
- Function gets lock before accessing globals; e.g., `malloc`

2. Cancellation

A thread may be canceled:

```
int pthread_cancel(pthread_t target)
```

Similar to sending a kill signal to UNIX process

A thread has:

- Cancel “state” – enabled or disabled
- Cancel “type” — asynchronous or deferred

Cancellation State

“Enabled” means thread can be canceled

“Disabled” means thread cannot be canceled

Cancelation Type

Assuming thread's cancelation state is "enabled" ...

- "Asynchronous" means: when canceled, thread is killed immediately
- "Deferred" means: thread may be canceled only at certain "cancelation points" where implementation checks "have I been canceled?" (and if canceled it kills itself)

Deferred cancelation is intended to help with the problem (which exists with signals at the process level) of a thread being canceled at "the worst moment" — leaving some data structure only partly updated, for example

Danger of Asynchronous Cancellation

Q: How to ensure that canceled thread won't be half-done with some operation that should not be left partially done (e.g., transfer of funds)?

A: Can't. Therefore, set thread cancel type to deferred before any sensitive operation:

1. set cancel type to deferred
2. perform sensitive operation to completion
3. set cancel type to asynchronous
4. test if thread was canceled:

```
if (pthread_testcancel())  
    pthread_cancel(pthread_self())
```

Cancelation Points

4 cancelation points in Pthreads implementation:

1. `pthread_testcancel`
2. `pthread_join`
3. `pthread_cond_wait`
4. `pthread_cond_timedwait`

POSIX states vendors *must* implement cancelation points in 23 specific library functions — roughly, those that may block

POSIX states vendors *may* implement cancelation points in approximately 50 other specific library functions

3. Adapting UNIX Semantics

Key issues:

- Process management
- Signals — to which thread should a signal be delivered?

Adapting UNIX Process Management

Examples of issues that arise when converting classic single-thread UNIX process model to modern multi-threaded processes ...

A. Does `fork` of N-thread process create another N-thread process or a 1-thread process?

B. Does `_exit(2)` terminate just one thread or whole process?

C. What happens to other threads when one thread calls `exec`?

A. Fork

When thread calls `fork` ...

- New 1-thread process is created
- Thread is replica of specific thread in parent process that called `fork`
- Address space of child duplicates that of parent — including all state created by other threads in parent (Yuck!)

Atfork

`pthread_atfork` function exists to help manage potential mess:

```
int pthread_atfork(void (*prepare)(void),  
                  void (*parent)(void),  
                  void (*child)(void));
```

prepare function called in parent before fork

parent function called in parent after fork

child function called in child after fork

B. Exit

Entire process terminates when any of these events occurs:

- Any thread calls `_exit(2)`
- Thread running `main` terminates
- Fatal signal is delivered to any thread

C. Exec

When any exec call happens:

- All existing threads terminated
- New thread created to run `main` of new executable file

Threads and Signals, I

POSIX added:

- Notion of per-thread signal mask
- Thread analogues of signal system calls:
`pthread_kill(pthread_t, int),`
`pthread_sigmask, etc.`

Each thread can mask signals individually

How thread handles signal depends on how signal was generated

Threads and Signals, II

If signal generated by hardware or software exception (e.g., SIGILL or SIGSEGV) ...

then “effective target” of signal is thread that caused exception, so ...

signal is delivered to offending thread

Threads and Signals, III

If signal generated by `pthread_kill` ...

then “effective target” of signal is specific thread, so ...

signal is delivered to targeted thread

Threads and Signals, IV

If signal generated by external process ...
(e.g., SIGCHLD or SIGUSR1) ...

then “effective target” is process, so ...

signal is delivered to *arbitrary thread that does not have signal blocked*

OS likely chooses the thread based on
what’s simplest to implement

4. Synchronizing Access to Shared Variables

Covered in next segment of the course ...