

Criticism

Use of lock for mutual exclusion satisfies all 5 criteria

However ...

- Depends on machine-specific instruction
- Looping (aka **busy-waiting** or **spinning**) by wait-er is wasteful of CPU time

Better solution might “suspend” waiter (i.e., place it on a queue) and “resume” it once lock available

- Locking doesn’t provide *coordination*—no way to tell waiting thread that it can now proceed because lock has been dropped

Mechanisms for Mutual Exclusion

1. Disable interrupts
2. *Without* hardware atomic instructions other than single-word read or single-word write
3. Hardware instructions for locking

WE ARE HERE

4. Hardware instructions for atomic access/update of certain common data structures
5. Higher level solutions (built on lower level solutions)

4. Compare-and-Swap

Compare-and-swap: compare memory location versus value held in register, then swap value in *second* register with memory location if compared values are equal

Two-word compare-and-swap: performs two comparisons, updates two locations if both comparisons are equal

CAS Semantics, I

Semantics of CAS(r1, r2, mem): atomically perform these actions ...

```
if (r1 == mem) {    /* compare */
    mem = r2        /* swap r2 & mem */
    r2 = mem
}
set Z CC bit based on (r1-mem)
```

CAS Semantics, II

Sometimes defined as CAS(r1, r2, mem) ...

```
if (r1 == mem) {    /* compare */
    mem = r2        /* not a true swap */
else
    r2 = mem
}
set Z CC bit based on (r1-mem)
```

Influential Motorola 680x0 implemented first
TAS & last CAS definitions

Expressive Power

CAS “more powerful” than TAS because of possibility of conditional action provided by comparison & 2nd register

2-word CAS “more powerful” than CAS because can update 2 locations

More powerful \Rightarrow can do more in a *single instruction*

(Can always get a lock then do what you want in several instructions)

Use of CAS for Locking

Assume: *first* definition of CAS

```
        mov r4 <- 0
loop:   mov r5 <- 1

//
// if (lock is 0)
//   then set it to 1
//   else do nothing
//
        cas r4, r5, lock

//
// if (lock was 1) try again;
// jump back to "mov" to be
// cautious -- lock's value
// (now in r5) should have been 1
// if other threads are
// well behaved but it doesn't
// hurt to be cautious
//
        bnz loop
```

Use of CAS for Increment

```
// read old value
loop:  mov r4 <- counter

// increment in r5
      mov r5 <- r4
      inc r5

//
// if (counter still equals value read)
//   then swap counter & new value
//
      cas r4, r5, counter

// if (counter != value read) try again
      bnz loop
```


Use of CAS for Insert into Singly-linked List

Assume:

- Within list element, next pointer is at offset NEXT_OFF
- Insert at head

CAS Insert

```
// r5 permanently points to new element
    mov r5 <- new

// r4 points to current head pointer
loop:  mov r4 <- head

// set new->next = current head
    mov (r5,NEXT_OFF) <- r4


//
// if (head is unchanged)
//   then swap to make head = new
//
    cas r4, r5, head

// if (head was changed), try again
    bnz loop
```

Illustration

After


```
    mov r5 <- new  
loop: mov r4 <- head
```



Illustration

After

```
mov (r5,NEXT_OFF) <- r4
```




Illustration

After

```
cas r4, r5, head
```

assuming $r4 = \text{head}$

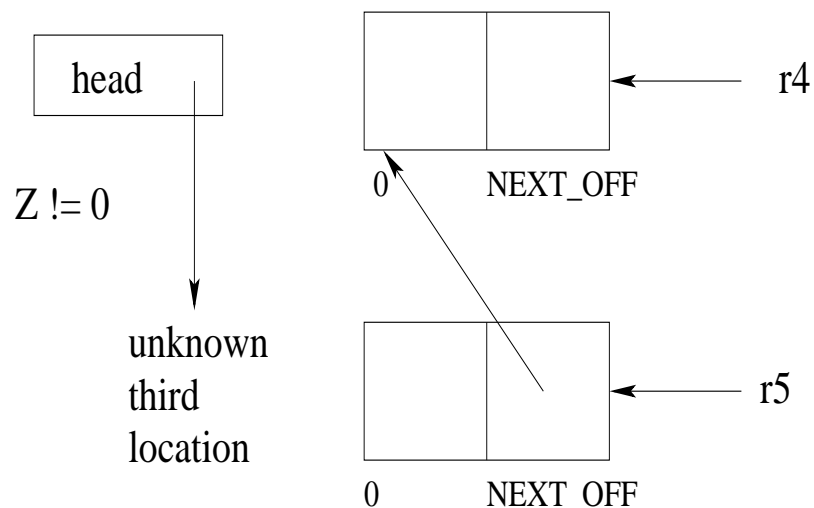


Illustration

After

```
cas r4, r5, head
```

assuming $r4 \neq head$



Summary of the Preceding Slides

TAS instruction was designed with only one use in mind – to get a lock atomically

CAS instruction was designed to get a lock atomically AND to do certain other common actions atomically

4. Load-Linked and Store-Conditional

TAS and CAS lock the processor-memory bus for their duration

- Bus is the bottleneck resource in many uniprocessor systems
- Even more so in many shared-memory multiprocessors

LL and SC designed with multiprocessing in mind

Assumes memory word can be “marked” with identity of marking CPU

Pioneered by MIPS architecture

LL and SC Semantics

Load Linked (LL) atomically does:

1. load memory location into register
2. mark location with CPU's ID

(Note: memory location can be marked by >1 CPUs)

Store Conditional (SC) atomically does:

1. if memory location is marked *by this CPU*, store new value and remove ALL marks; else do nothing
2. return indication of which case occurred

Use of LL+SC for Locking


To get lock:

```
r2 = 1;
while (1) {
    LL r1, lock
    if (r1 == 0)
        if (SC r2, lock)
            break;
}
```

SC returns 1 iff lock was STILL marked by this CPU; i.e., no other CPU executed SC before this one did

To drop lock: STORE lock, 0

Example: No Contention




No marks

0

 lock

LL r1, lock




Marked

0

 lock

SC r2, lock




No marks


1

 lock


Example: Contention




CPU A: LL r1, lock



CPU B: LL r1, lock



CPU A: SC r2, lock



CPU B: SC r2, lock ... do nothing because there is no B mark;
loop and perform LL and SC instructions again

CPU B: LL r1, lock ... r1 is 1 so keep looping

Use of LL+SC for Increment

```
while (1) {  
    LL r1, counter  
    r2 = r1 + 1;  
    if (SC r2, counter)  
        break;  
}
```

Summary

Special atomic instructions can be used to implement locks AND ALSO to perform other common operations

Which other operations depends on the instruction

But: these “fancy uses” are rarely employed

Mechanisms for Mutual Exclusion

1. Disable interrupts
2. *Without* hardware atomic instructions other than single-word read or single-word write
3. Hardware instructions for locking
4. Hardware instructions for atomic access/update of certain common data structures

WE ARE HERE

5. Higher level solutions (built on lower level solutions)

Pthreads Locking API

`pthread_mutex_lock` — waits if lock already held

`pthread_mutex_trylock` — immediately returns indication whether lock is already held

`pthread_mutex_unlock` — drops lock

Mutex Concept

Q: What is a pthread “mutex” lock?

A: A simple lock – either it’s held or it’s not, nothing fancier

Commonly implemented with:

- special instructions mentioned above
- semaphore initialized to 1 (and semaphore is implemented using special instructions)