

# Pthreads Condition Primitives

Pthreads API for condition variables:

```
pthread_cond_init(pthread_cond_t *c,  
                  pthread_cond_attr_t *attr)
```

```
pthread_cond_destroy(pthread_cond_t *c)
```

```
pthread_cond_wait(pthread_cond_t *c, pthread_mutex_t *m)
```

```
pthread_cond_signal(pthread_cond_t *c)
```

```
pthread_cond_broadcast(pthread_cond_t *c)
```

# Use of Pthreads

## Mesa-style Condition Variables, I

Operation of `pthread_cond_wait(c, m)`:

1. Set thread state to “waiting”
2. Add thread to set of threads waiting for condition variable `c`
3. `pthread_mutex_unlock(m)`
4. Some time later: thread state set to “ready” by another thread’s signal
5. Eventually, scheduler will pick this thread to run; then its state will be “running”
6. `pthread_mutex_lock(m)`

Operation of `pthread_cond_signal(c)`:

- Select thread from `c`’s wait-set according to policy & set its state to “ready”

# Use of Pthreads

## Mesa-style Condition Variables, II

```
pthread_mutex_t mutex;          // explicit monitor lock

pthread_cond_t spaces, items;   // condition variables

// this is a monitor procedure
void produce() {
    pthread_mutex_lock(&mutex);   // get monitor lock
    while (<there is no space>)    // first action:
        pthread_cond_wait(&spaces, &mutex); // verify state

    ... produce ...              // change state

    pthread_cond_signal(&items);  // last action:
                                   signal state change
    pthread_mutex_unlock(&mutex); // drop monitor lock

    return;
}
```

# Use of Pthreads

## Mesa-style Condition Variables, III

Mutex implemented by programmer, not language, enforces monitor invariant

Data must be in state associated with app-specific condition “spaces” BEFORE produce operation runs, and will be in state associated with app-specific condition “items” AFTER operation runs

# Use of Pthreads

## Mesa-style Condition Variables, IV

In Hoare's formulation, after `signal()`, signaled thread is guaranteed next to run

Therefore, wait with:

```
if (<not condition>)  
    wait();
```

# Use of Pthreads

## Mesa-style Condition Variables, V

In Pthreads for C/UNIX ...

1. Other threads may run between signal-er and signal-ee

Between when thread A's `signal()` awakens thread B and when thread B waits for the lock, some thread C may have run and undone the application-specific condition that A established for B

# Use of Pthreads

## Mesa-style Condition Variables, VI

2. Therefore, condition must be re-checked when signal-ee awakens

3. Therefore, waiters wait with:

```
while (<not condition>)  
    wait();
```

# Example, I

Ideal execution:

0. Thread A holds lock & is in monitor procedure
1. Thread A establishes application-specific condition on monitor data
2. Thread A signals appropriate condition variable, thereby changing state of Thread B from waiting to ready
3. Thread A drops lock
4. Thread A returns from monitor procedure
5. Thread B (which is in `pthread_cond_wait`) is scheduled
6. Thread B tries to get lock
7. Thread B succeeds in getting lock & its `pthread_cond_wait` returns
8. Thread B tests application-specific condition in test of while statement
9. Test is passed & thread B enters monitor procedure



# Example, II

In Mesa formulation (OS thread scheduler doesn't know about monitors), what COULD happen:

... same as above, steps 0 thru 4

4.25. Thread C is scheduled

4.50. Thread C gets lock, enters a monitor procedure  
and changes monitor data so as to UNDO thread B's  
application-specific condition

4.75. Thread C drops lock & leaves its monitor procedure

5. Thread B (which is in pthread\_cond\_wait) is scheduled

6. Thread B tries to get lock

7. Thread B succeeds in getting lock & its pthread\_cond\_wait  
returns

8. Thread B tests application-specific condition in test of  
while statement

9. Test FAILS!

10. Thread B loops & calls pthread\_cond\_wait again

Intervention by thread C shows why thread  
B must re-test application-specific condition  
after returning from wait

# Starvation

In example above, for thread B to eventually enter monitor procedure all these must happen:

- a. Some thread A gets into monitor
- b. Thread A establishes B's application-specific condition
- c. Thread A signals condition B is waiting for
- d. B is awakened **WITHOUT** an intervening thread undoing the condition as Thread C did in example above

There is **NO** guarantee that all this will eventually happen—B may be starved

(E.g., suppose d never happens ... there is always some intervening thread that undoes condition)

# Use of Pthreads

## Mesa-style Condition Variables, VII

Above *conventional* use of wait and signal provides proper monitor entry & exit

Costs are:

1. Extra evaluation of condition (after return from wait)
2. Possible starvation, since no scheduler cooperation

Benefits are:

1. Decouples OS thread scheduler policy from monitors—can have monitors without language or OS support
2. Does not require forced (and possibly non-optimal) Hoare-style immediate switch to signaled thread

# How to Program with Monitors, I

1. Identify logically related variables that might be accessed concurrently & place them inside monitor
2. Determine possible states for these variables
3. Determine state transitions that might occur concurrently
4. Write each state transition as a monitor procedure

# How to Program with Monitors, II

Data can now be transitioned from one state to the next safely

Each monitor procedure performs a state transition following this pattern:

```
while (data not in state A)
    wait(cv for state A)
<transition data from state A to state B>
signal(cv for state B)
```

# Signal vs. Broadcast, I

When to use broadcast?

`pthread_cond_broadcast` man page says:

The `pthread_cond_broadcast()` function is used whenever the shared-variable state has been changed in a way that more than one thread can proceed with its task.

Consider a single producer/multiple consumer problem, where the producer can insert multiple items on a list that is accessed one item at a time by the consumers.

By calling the `pthread_cond_broadcast()` function, the producer would notify all consumers that might be waiting, and thereby the application would receive more throughput on a multiprocessor.

# Signal vs. Broadcast, II

Signal vs. broadcast is only a performance optimization

For code to be correct, it must ALWAYS be possible to change any signal to a broadcast (or vice versa)