

# Monitor-like Behavior

- Java provides language support for monitor lock: **synchronized** keyword
- At most one of foo, bar may run at any moment:

```
public class ThreadExample implements Runnable {  
    private int var1;    // instance variables  
    private int var2;  
  
    public synchronized void foo() {  
        // statements accessing var1 and/or var2  
    }  
  
    public synchronized void bar() {  
        // statements accessing var1 and/or var2  
    }  
}
```

# Synchronized Keyword

- synchronized means: at most one of all the methods declared synchronized may be executing at a time
- This applies across ALL threads
- This is not the complete monitor concept—that requires “wait” and “signal” operations—but it is the “monitor lock” part of the concept

# Object Lock

- Each object has one hidden **object lock** that is ...
  - Obtained upon entry to synchronized method
  - Dropped upon exit from synchronized method
- Lock is dropped no matter how method exits, e.g.,
  - By return
  - By throwing exception
  - By failing to catch thrown exception

# Object Lock is Smart

It is OK to multiply acquire the lock:

```
public class ThreadExample implements Runnable {  
    public synchronized void foo() {  
        bar();  
    }  
  
    public synchronized void bar() {  
        ...  
    }  
}
```

# Smart Does Not Mean Idiot-Proof

- Suppose programmer declares `run` to be synchronized:

```
public class ThreadExample implements Runnable {  
    public synchronized void run() {  
        bar();  
    }  
  
    public synchronized void bar() {  
        ...  
    }  
}
```

- Since thread is ALWAYS executing `run`, other threads could never invoke `bar`

# Synchronized and Unsynchronized Methods, I

- Consider:

```
public class ThreadExample implements Runnable {  
    public synchronized void foo() {  
        bar();  
    }  
  
    public void bar() {  
        ...  
    }  
}
```

- ANY number of threads may simultaneously invoke bar
- Only ONE thread may invoke foo

Q: What happens if foo calls bar while N other threads are executing bar?

# Synchronized and Unsynchronized Methods, II

A: Thread enters bar – so that  $N+1$  threads are executing bar

- All “synchronized” means is: “get (reentrant) object lock when entering this method & drop it when exiting”
- Even though foo is synchronized there is no restriction on bar

# Synchronized Blocks, I

- synchronized keyword can be applied to units smaller than entire method
- This is a performance optimization
- Consider:

```
public synchronized void silly() {  
    1. ... compute Pi to a billion digits ...  
    2. ... store result in shared memory ...  
}
```

- Need lock during #2 (very short) but not during #1 (very long)



# Synchronized Blocks, II

- Can write as:

```
public void silly() {  
    1. ... compute Pi to a billion digits ...  
    synchronized(this) {  
        2. ... store result in shared memory ...  
    }  
}
```

- Lock held only during #2
- Argument “(this)” specifies which object’s lock to obtain
- Recall: `this` refers to current object

# Synchronized Static Methods

Q: Can static methods be synchronized?

A: Yes

Q: How?

A: Each CLASS has one hidden **class lock**

- Entry into synchronized static method gets class lock
- Therefore: at most one of class's static methods declared synchronized may be executing at a time

# Java Memory Model

- Java has a **memory model**

- See

`www.cs.umd.edu/~pugh/java/memoryModel/jsr-133-faq.html`

- Memory model specifies required behavior of underlying memory (as implemented by hardware and JVM)
- Mostly relevant to caching & multiprocessors, esp. NUMAs
- One aspect of Java memory model: loads & stores of any one-word variable must be atomic (except long and double, which are each 64 bits)

# Another Way to Synchronize: Volatile

- You can declare an ATOMIC variable to be “volatile”
- volatile variable will not be cached, meaning:
  - Current value always read from memory
  - New value always written to memory
- volatile does not apply to non-atomic types: long, double

# Volatile:

## What It's Good For

- If one method assigns to int and another reads int, can declare int as volatile and avoid using synchronized
- Use of volatile makes this class thread-safe without “synchronized” keyword:

```
public class ThreadExample implements Runnable {  
    private volatile int x;  
  
    public void set(int arg) {  
        x = arg;    // atomic write to int x  
    }  
  
    public int get() {  
        return x;    // atomic read of int x  
    }  
}
```

# Volatile:

## What It's Bad For, I

- My opinion: `volatile` is dangerous – don't use it!
- Very easy to misuse
- For example, this code is **WRONG** – does not provide synchronized access:

```
private volatile int foo;  // declaration
...
foo++;                    // use in some method
```

Q: Do you understand why?

# Volatile:

## What It's Bad For, II

A: “foo++” compiles to instructions that load AND store foo, though it is easy to think of it as only a store

- This also is wrong:

```
private volatile int[] foo;  
    ...  
foo[4] = 17;
```

- Reason: “foo” alone — the REFERENCE to the array — is volatile
- Individual array elements are NOT made volatile by the declaration

## Aside: Volatile Keyword

- C, C++, Java, and C# all have the “volatile” keyword
- Keyword has slightly different meaning in all languages!
- More danger: you may think you know what volatile means but you are mistaken