

## Aside: Volatile Keyword

- C, C++, Java, and C# all have the “volatile” keyword
- Keyword has slightly different meaning in all languages!
- More danger: you may think you know what volatile means but you are mistaken

## Aside: History of Volatile

- C introduced `volatile` “type qualifier” keyword
- Meant to accommodate memory mapped devices
- Keyword warns that variable’s value “can be changed by means outside this code”

# Example, I

```
/* int foo declared mapped to hardware device register */

void read() {
    foo = 0;

    while (foo != 1)
        ;    /* spin-wait for foo to become 1 */

    ... code that assumes foo == 1 ...
}
```

Q: What's going on in the loop?

A: foo represents a hardware device register whose value might be changed asynchronously by the hardware

## Example, II

- Optimizing compiler will emit same assembly code as for:

```
void read() {  
    foo = 0;  
  
    while (1)  
        ;    /* infinite loop */  
  
    ...  
}
```

- Loops forever without ever testing foo's value

## Example, III

- Declaring `foo` to be `volatile` fixes the problem:

```
static volatile int foo;

void read() {
    foo = 0;

    while (foo != 1)
        ;    /* spin-wait for foo to become 1 */

    ... code that assumes foo == 1 ...
}
```

- `volatile` keyword tells compiler always to emit instructions to read/write `foo`'s memory location

# Summary of Volatile

- `volatile` keyword is:
  - Hard to use correctly
  - Easy to use incorrectly
  - Subtly different in different languages
  - Unnecessary (for Java application programming)
- Summary: avoid `volatile` unless you really understand what you're doing and you're sure there is no other way to do it

# Library Support for Concurrency

1. Atomic classes
2. Semaphore, read/write locks, Barrier, Countdown latch, Exchanger, SynchronousQueue
3. Collection classes, concurrent data structures
4. Task management (thread pools)
5. Concurrency patterns

# Summary

- Java = base language +  
HUGE library
- Slogan for conceptual cleanliness ...  
“less is more”
- Slogan for Java language/library  
designers??? ... “more is more”
- In their defense: why NOT accumulate  
debugged classes?
- But ... library is in serious need of an  
intelligent table of contents; i.e., advice for  
which library classes to use, when, and how



# Reminder:

## Wrapper Classes

- Java has 8 primitives types: `int`, `long`, `boolean`, etc.
- Recall that each primitive type has a corresponding *wrapper class*
- E.g., `Integer` for `int`, `Boolean` for `boolean`, etc.
- Allows treating primitive variable as an object
- Houses useful static methods; e.g., `Integer.parseInt`

# Atomic Classes, I

- 3 of 8 primitive types (`int`, `long`, `boolean`) have corresponding **atomic classes**
- They are `AtomicInteger`, `AtomicLong`, `AtomicBoolean`
- Others too ... `AtomicReferenceArray` etc.
- These are part of `java.util.concurrent.atomic`
- Methods of atomic classes each behave atomically

# Atomic Classes, II

- All atomic classes have methods `get`, `set`, `getAndSet`, `compareAndSet`
- `compareAndSet(expect, update)` sets new value ( “update” ) only if old value matches expected value ( “expect” )
- Just like test-and-set machine instruction
- `AtomicInteger` and `AtomicLong` also have

<code>incrementAndGet</code>	<code>++x</code>
<code>getAndIncrement</code>	<code>x++</code>
<code>decrementAndGet</code>	<code>--x</code>
<code>getAndDecrement</code>	<code>x--</code>

- Other methods as well: `addAndGet`, etc.

# Utility of Atomic Classes

1. Simple compound operations that cannot be done with a `volatile` variable

E.g., `x++` requires a read AND a write

2. Building blocks for more complicated classes

# Atomic Array Classes

- `AtomicIntegerArray`, `AtomicLongArray`, `AtomicReferenceArray` provide atomic access to elements of array
- Other classes as well

# Other Synchronization Classes

- Semaphore — we've seen this before
- Read/write locks — we've seen this before
- Barrier — we've seen this in Assignment 3
- Countdown latch — specialization of barrier
- Exchanger — even more specialized
- SynchronousQueue — not really a queue
- Locks are part of  
`java.util.concurrent.locks`,  
rest are part of `java.util.concurrent`

# Semaphore, I

- Initial value of semaphore called “number of permits:”

`Semaphore(long permits)`

`Semaphore(long permits, boolean fair)`

- “fair” access means FCFS — goal is to avoid starvation

# Semaphore, II

- P and V called “acquire” and “release:”

`acquire()` throws `InterruptedException`

`acquireUninterruptibly()`

`acquire(long permits)` throws `InterruptedException`

`acquireUninterruptibly(long permits)`

`release()`

`release(long permits)`

- “permits” argument adds/deletes N to semaphore value
- There are various “tryAcquire” methods as well



# Interruptibility

- You will see methods named `“fooInterruptibly”` and `“fooUninterruptibly”`
- Must be aware of what the default behavior is

# Reminder: Moded Locks

- Recall: purpose of adding **modes** to lock is to increase concurrency
- Separate threads can simultaneously hold lock in “compatible” modes; e.g., read-read
- BY FAR most common set of modes is read & write
- There are other “fancy” mode sets
- (SIX mode is very elaborate and useful only for hierarchical data)

# Read/Write Locks

Interface:

```
public interface ReadWriteLock {  
    Lock readLock();  
    Lock writeLock();  
}
```

Class:

```
public class ReentrantReadWriteLock  
    implements ReadWriteLock {  
    public ReentrantReadWriteLock();  
    public ReentrantReadWriteLock(boolean fair);  
    public Lock readLock();  
    public Lock writeLock();  
}
```

# Read/Write Lock

## Semantics, I

- Reentrant: lock can be multiply acquired by same thread
- OK to downgrade, like so:
  1. Acquire write lock
  2. Acquire read lock
  3. Drop write lock
- Upgrade NOT OK — must first drop read lock, then try to acquire write lock

# Read/Write Lock Semantics, II


- Unusual API
- Uses **nested classes** of  
ReentrantReadWriteLock:  
ReentrantReadWriteLock.ReadLock and  
ReentrantReadWriteLock.WriteLock
- How to perform the 4 operations:

```
ReentrantReadWriteLock rwl = new ReentrantReadWriteLock();  
rwl.readLock().lock();  
rwl.readLock().unlock();  
rwl.writeLock().lock();  
rwl.writeLock().unlock();
```

# Barrier


- **Barrier** is common synchronization primitive
- ALL threads must reach barrier before ANY thread may continue past barrier
- Most often used in scientific computation; e.g., relaxation computation

# Depiction



## Aside: Relaxation, I

- Common pattern in scientific computation: simulate 2D or 3D real-world phenomenon by computing the physics at each point on a grid:




- The denser the grid points, the more faithful the simulation



## Aside: Relaxation, II

- Values at adjacent grid points affect the computation
- Simulation runs in stages: compute result at each point, exchange results, compute again, etc.



- End of each stage is a barrier synchronization point

# Barrier API

```
public class CyclicBarrier {  
    public CyclicBarrier(int parties);  
    public CyclicBarrier(int parties,  
                           Runnable barrierAction);  
    public int await() throws InterruptedException,  
                           BrokenBarrierException;  
    public boolean isBroken();  
    public void reset();  
    public int getParties();  
    public int getNumberWaiting();  
}
```

# Barrier Use, I

1. Each thread calls `await()` and blocks
  2. Once LAST thread calls `await()`, ALL threads return
- Barrier can be reused in this fashion — that's why it's called "*cyclic*"
  - The number of threads given as "parties" constructor argument

## Barrier Use, II

- Possible to specify action to take after all threads reach barrier but before they are released:

```
public CyclicBarrier(int parties,  
                    Runnable barrierAction);
```

- “barrierAction” is an object that implements Runnable — meaning it has a “run” method
- run() is executed by last thread to reach barrier (NOT in a separate thread)

# Broken Barrier Concept

- Barrier may be “broken:”

```
public int await() throws InterruptedException,  
                    BrokenBarrierException;  
public boolean isBroken();
```

- Ways to break a barrier:

- await-ing thread is interrupted
- barrierAction throws exception

- Broken barrier must be “reset” before next use:

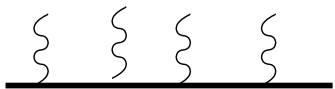
```
reset()
```

# Problems With Barriers

1. Must know number of threads when barrier is created
2. Scheduling anomaly can lead to barrier getting stuck

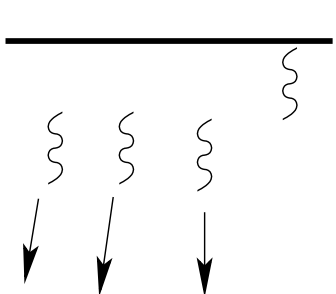
# Scheduling Anomaly

Barrier often occurs inside loop; e.g.,  
relaxation problem



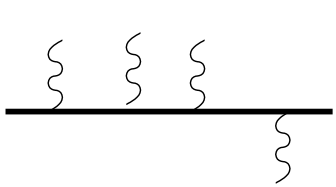
all threads at barrier

The diagram shows a horizontal line representing a barrier. Four curly braces are positioned above the line, each representing a thread waiting at the barrier.



all threads released  
from barrier, but  
one doesn't get  
scheduled promptly

The diagram shows a horizontal line representing a barrier. Three curly braces are positioned above the line, each with a downward-pointing arrow below it, indicating they have been released. A fourth curly brace is positioned to the right of the line, indicating a thread that has not yet been released.



at next occurrence  
of barrier, all threads  
are "at the barrier"  
but one thread did  
not perform its  
computation during  
the previous stage

The diagram shows a horizontal line representing a barrier. Three curly braces are positioned above the line, each representing a thread that has completed its computation. A fourth curly brace is positioned to the right of the line, representing a thread that has not yet completed its computation.

# Solution

- Need a **sense reversing barrier**
- Java's `CyclicBarrier` is sense reversing – that's what “cyclic” means



# Non-Cyclic Barrier

- Non-cyclic barrier's internal logic executed by each thread:

```
LOCK(barrier.lock);
if (barrier.count == 0) {
    // first thread executes here
    barrier.release = 0;
}
local_count = ++barrier.count;
UNLOCK(barrier.lock);
if (local_count == N) {
    // last thread executes here
    barrier.count = 0;
    barrier.release = 1;
} else {
    while (barrier.release == 0)
        ;
}
```

- Unprotected access to barrier inside if-clause is OK because thread knows others are waiting (spinning on barrier.release)
- Unprotected access to barrier.release assumes it is atomic (i.e., one word)

# Problem With Non-Cyclic Barrier

- One “fast” thread ...
  - Releases from first barrier
  - Is first thread to get to second barrier
  - Resets `barrier.release` to 0

All this happens before “slow” thread releases from first barrier

- Slow thread is still executing

```
while (barrier.release == 0)
    ;
```

- Slow thread won't release from first barrier so count at second barrier will never go beyond N-1

# Cyclic Barrier

Cyclic barrier's internal logic executed by each thread:

```
local_sense = !local_sense;
LOCK(barrier.lock);
local_count = barrier.count++;
UNLOCK(barrier.lock);

if (local_count == N) {
    barrier.count = 0;
    barrier.release = local_sense;
} else {
    while (barrier.release != local_sense)
        ;
}
```

# Countdown Latch

- Similar to barrier
- `CountDownLatch` starts with a positive int value
- Each thread to call `await()` blocks if `count > 0`
- When count reaches 0, all threads released

Q: How does count go down?

A: call "`countDown()`"

# Exchanger

- Think of it as 2-thread barrier plus object exchange:

```
public class Exchanger<V> {  
    public Exchanger();  
    public V exchange(V x) throws  
        InterruptedException;  
}
```

- A *generic* type — note type parameter “V”
  1. First thread to call “exchange” blocks
  2. When second thread calls exchange, both threads return AND each returns the other’s argument

# SynchronousQueue

- Think of it as 1-item exchanger
- Implements what textbooks call a “rendezvous”
- First thread calls `put` and blocks
- Second thread later calls `take` – it now has the item and the first thread is unblocked
- (If `take` is called first then `put` unblocks the taker)
- Despite the name, it's NOT a queue ... `size` always returns 0

# Does This Stuff Really Get Used?

Semaphore — yes

Read/write locks — yes

Barrier — yes

Countdown latch — no

Exchanger — no

SynchronousQueue — no